


A vibrant photograph of an orange grove. In the foreground, two rows of orange trees are filled with ripe, bright orange fruit and lush green leaves. A dirt path leads straight down the center of the grove, flanked by the trees. In the background, a range of blue mountains is visible under a clear blue sky with a few wispy white clouds.

Orange Analysis

Michael Brown - Jorge Terrazas

Objectives

- 🍊 Determine the most important attributes in determining an orange's quality.
- 🍊 Predict the quality of an orange from its measurable features using a regression model
- 🍊 Identify orange quality clusters using a clustering model





Context: Real World Applications

- 🍊 Uphold food quality standards
- 🍊 Reduce unnecessary food waste
- 🍊 Identify trends for pricing



About the Source Data

-  Numerical attributes describing the quality of oranges, including their size, weight, sweetness (Brix), acidity (pH), softness, harvest time, and ripeness.
-  Categorical attributes such as color, variety, presence of blemishes, and overall quality.

Data Wrangling

Cleanup, preprocessing

dataset:

```
OrangeDF.head(5)
```

✓ 0.0s

	Size (cm)	Weight (g)	Brix (Sweetness)	pH (Acidity)	Softness (1-5)	HarvestTime (days)	Ripeness (1-5)	Color	Variety	Blemishes (Y/N)	Quality (1-5)
0	7.5	180	12.0	3.2	2.0	10	4.0	Orange	Valencia	N	4.0
1	8.2	220	10.5	3.4	3.0	14	4.5	Deep Orange	Navel	N	4.5
2	6.8	150	14.0	3.0	1.0	7	5.0	Light Orange	Cara Cara	N	5.0
3	9.0	250	8.5	3.8	4.0	21	3.5	Orange-Red	Blood Orange	N	3.5
4	8.5	210	11.5	3.3	2.5	12	5.0	Orange	Hamlin	Y (Minor)	4.5

preprocessing & cleaning data:

```
print(OrangeDF["Color"].apply(type).unique())  
print(OrangeDF["Variety"].apply(type).unique())  
print(OrangeDF["Blemishes (Y/N)"].apply(type).unique())  
print(OrangeDF.isna().sum())
```

[5] ✓ 0.0s

Python

```
... [<class 'str'>]  
[<class 'str'>]  
[<class 'str'>]  
Size (cm)      0  
Weight (g)     0  
Brix (Sweetness) 0  
pH (Acidity)   0  
Softness (1-5) 0  
HarvestTime (days) 0  
Ripeness (1-5) 0  
Color          0  
Variety        0  
Blemishes (Y/N) 0  
Quality (1-5)  0  
dtype: int64
```

checking for missing or wrong format values:

^No columns with mixed values and no missing values in dataset


Data Analysis

Distributions, basic relationships


checking for outliers & data distribution of columns:

	Size (cm)	Weight (g)	Brix (Sweetness)	pH (Acidity)	Softness (1-5)	HarvestTime (days)	Ripeness (1-5)	Quality (1-5)
count	241.000000	241.000000	241.000000	241.000000	241.000000	241.000000	241.000000	241.000000
mean	7.844813	205.128631	10.907884	3.473900	3.072614	15.344398	3.599585	3.817427
std	1.086002	56.461012	2.760446	0.421007	1.323630	5.323852	1.205214	1.014410
min	6.000000	100.000000	5.500000	2.800000	1.000000	4.000000	1.000000	1.000000
25%	6.900000	155.000000	8.500000	3.200000	2.000000	11.000000	3.000000	3.000000
50%	7.800000	205.000000	11.000000	3.400000	3.000000	15.000000	4.000000	4.000000
75%	8.700000	252.000000	13.400000	3.800000	4.000000	20.000000	4.500000	4.500000
max	10.000000	300.000000	16.000000	4.400000	5.000000	25.000000	5.000000	5.000000


No extreme outliers and fairly normally distributed data for most columns as shown below




correlation
matrix to
check for
correlations
between
features

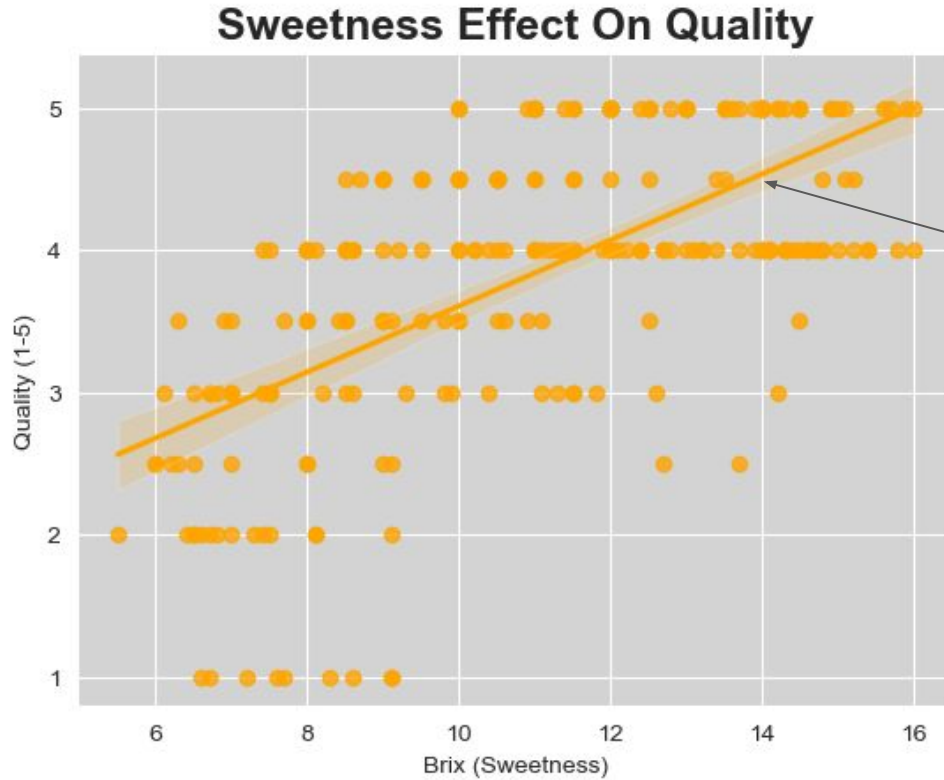


Highest
correlation
is between
Sweetness
and Quality



**Largest
negative
correlation
is between
Harvest
Time and
Quality**



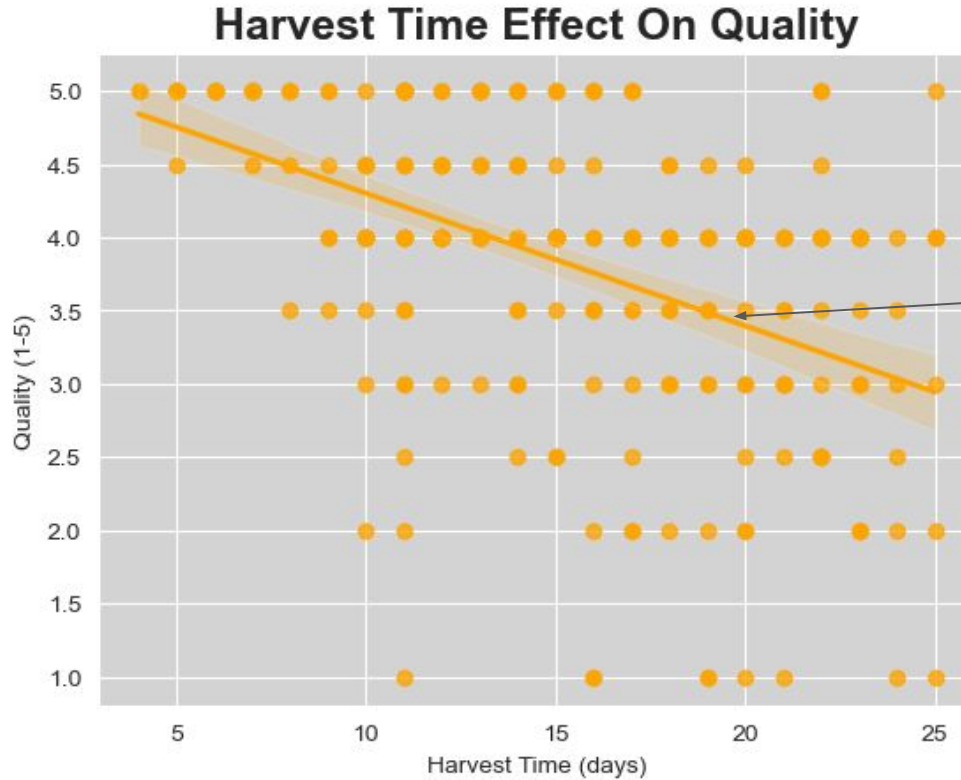


Visual relationship
between:
Sweetness & Quality

Relatively strong *positive*
relationship

What does this tell us?

Visual relationship between: Harvest Time & Quality



Relatively strong *negative* relationship

Implication?

Models

Baseline, regression, clustering

Baseline Model

```
baseline = OrangeDF["Quality (1-5)"].mean()
mean_guesses = np.full(OrangeDF.shape[0], baseline)
print("mean baseline guess r^2:", r2_score(OrangeDF["Quality (1-5)"], mean_guesses))
print("mean baseline guess MAE:", mean_absolute_error(OrangeDF["Quality (1-5)"], mean_guesses))
print("mean baseline guess MSE:", mean_squared_error(OrangeDF["Quality (1-5)"], mean_guesses))
```

✓ 0.0s

```
mean baseline guess r^2: 0.0
mean baseline guess MAE: 0.7897419121571596
mean baseline guess MSE: 1.0247585268848678
```

Quality Predictions with Regression Models

```
X = OrangeDF.drop(columns=['Quality (1-5)', 'Variety'])
y = OrangeDF['Quality (1-5)']

lin_reg = LinearRegression()
lin_scores = cross_validate(estimator=lin_reg, X=X, y=y, scoring = ('r2', 'neg_mean_absolute_error', 'neg_mean_squared_error'))
mlp_model = MLPRegressor(hidden_layer_sizes= 1000, activation="relu", early_stopping=True, max_iter=500)
mlp_scores = cross_validate(estimator=mlp_model, X=X, y=y, scoring = ('r2', 'neg_mean_absolute_error', 'neg_mean_squared_error'))
tree_model = DecisionTreeRegressor()
tree_scores = cross_validate(estimator=tree_model, X=X, y=y, scoring = ('r2', 'neg_mean_absolute_error', 'neg_mean_squared_error'))
print("LINEAR REGRESSION:")
print("r^2s:", lin_scores["test_r2"], "\nMAE:", -lin_scores["test_neg_mean_absolute_error"], "\nMSE", -lin_scores["test_neg_mean_squared_error"])
print("\nMULTILAYER PERCEPTRON:")
print("r^2s:", mlp_scores["test_r2"], "\nMAE:", -mlp_scores["test_neg_mean_absolute_error"], "\nMSE", -mlp_scores["test_neg_mean_squared_error"])
print("\nDECISION TREE:")
print("r^2s:", tree_scores["test_r2"], "\nMAE:", -tree_scores["test_neg_mean_absolute_error"], "\nMSE", -tree_scores["test_neg_mean_squared_error"])]
```

✓ 3.3s

Python

LINEAR REGRESSION:

r^2s: [0.39824754 0.42717193 0.23199958 0.23614747 0.22132411]

MAE: [0.63176573 0.53574796 0.72742815 0.54049215 0.69251087]

MSE [0.63496037 0.39375714 0.80391711 0.56286052 0.8922328]

MULTILAYER PERCEPTRON:

r^2s: [0.3950977 0.39415192 0.31429317 0.05016218 0.33410057]

MAE: [0.64016155 0.55822475 0.69835993 0.64626905 0.6906529]

MSE [0.63828404 0.41645481 0.71777493 0.69990762 0.76300977]

DECISION TREE:

r^2s: [0.54549043 0.74238358 -0.09961646 -0.43484023 -0.04545455]

MAE: [0.32653061 0.20833333 0.82291667 0.63541667 0.72916667]

MSE [0.47959184 0.17708333 1.15104167 1.05729167 1.19791667]

K-Means Clustering Model: Good or Moderate Quality

```
n_clusters = 2
KM = KMeans(n_clusters = n_clusters)
to_cluster = X[["Brix (Sweetness)", "HarvestTime (days)"]]
KM.fit(to_cluster)
labels = KM.labels_
cluster_df = pd.DataFrame({"quality score": y.to_list(), "cluster label": labels})
cluster_df.sort_values(by="cluster label")
for i in range(n_clusters):
    print("cluster", i, "median quality score:", cluster_df.loc[cluster_df["cluster label"] == i]["quality score"].median())
```

✓ 0.1s


cluster 0 median quality score: 3.5

cluster 1 median quality score: 4.5

Orange Variety Similarity Dendrogram

Distances based on:

- Sweetness (Brix)
- Acidity (pH)



Thank You

Citations

Source Data: <https://www.kaggle.com/datasets/shruthiiee/orange-quality/data>

<https://scikit-learn.org/stable/>

<https://matplotlib.org/>

<https://pandas.pydata.org/>

<https://numpy.org/>

<https://seaborn.pydata.org/>

<https://github.com/kiat/Elements-of-Data-Analytics>